

An Introduction to Reinforcement Learning



Farnaz Adib Yaghmaie

Linköping University, Sweden
farnaz.adib.yaghmaie@liu.se

March 12, 2021

A few words about the RL workshop

- Given by Farnaz Adib Yaghmaie and Lennart Ljung.
 - Fredrik Ljungberg will assist me today.
- The workshop is a part of LINK-SIC.
- We define the basic concepts in RL: a bit of theory is involved.
- The presentation is open to discussion:
 - Write your question in chat
- Run the code to understand better.

Machine Learning

- Supervised Learning
- Unsupervised Learning
- **Reinforcement Learning**

Finding suitable actions to take in a given situation in order to maximize a reward¹.

¹Richard S Sutton & Andrew G Barto. *Reinforcement learning: An introduction*, volume 1. MIT press Cambridge, 1998.



How RL is different from other branches of ML?

- No supervisor; only a reward
- The action will effect subsequent data
- Dynamic data vs. Static data

Examples of RL

- Maneuvering helicopter
- Defeating Backgammon champion
- Surpassing human performance in playing many Atari games
- Making a humanoid robot walk
- Research and find many others!
- Find a successful application in your field!

Beating Atari champions ²



² Mnih et al. *Playing Atari with Deep Reinforcement Learning*, arXiv preprint, 2013.

A graphical representation





Photo Credit: @ https://en.wikipedia.org/wiki/Reinforcement_learning

Three main components in RL

- Reward
- Environment
- Agent



Reward r_t

- Reward is a scalar signal
- It indicates how well agent is doing at time t
- The aim is to maximize the total reward
- Or equivalently minimize a total penalty

Environment

- Generates dynamic data
 - Receives the action a_t
 - Emits a new state s_{t+1}
- Has unknown rules; i.e. dynamics
- Emits a reward r_t

Agent

- Receives the reward r_t
- Receives the observation s_t
- Execute an action a_t

Generate actions to maximize the future rewards

Three main components of an RL agent

- Policy: The agent's rule to select action in a given state
- Value function: How good the agent's rule is
- Model of environment: The agent's interpretation of the environment

Not all are necessary!

Policy Gradient

Learning policy

Dynamic Programming based

Learning value function

Model building

Learning the model of environment

Let's play a game

You are playing chess with a Grandmaster of chess, Dorsa.
What are

- Environment?
- Reward?
- Agent?




Photo Credit: @ <https://www.pinterest.se/>

Email your questions to

farnaz.adib.yaghmaie@liu.se