

Markov Decision Process



Farnaz Adib Yaghmaie

Linköping University, *Sweden*
farnaz.adib.yaghmaie@liu.se


March 12, 2021

Markov Decision Processes

- describe environment in RL framework
- describe dynamical systems
- In optimal control problems MDPs are continuous

A Markov Decision Process (MDP) is a tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- \mathcal{S} : The set of states.
- \mathcal{A} : The set of actions.
- \mathcal{P} : The set of transition probability.
- \mathcal{R} : The set of immediate rewards associated with the state-action pairs.
- $0 \leq \gamma \leq 1$: Discount factor.



Modified version of @ https://en.wikipedia.org/wiki/Markov_decision_process

States: Describe internal status of MDP

Actions: Possible choices to make in each state of MDP

Transitions probability: \mathcal{P} is the set of transition probability with n_a matrices each of dimension $n_s \times n_s$ where s, s' entry reads

$$[\mathcal{P}^a]_{ss'} = p[s_{t+1} = s' | s_t = s, a_t = a] \quad (1)$$

Reward:

$$r_t = r(s, a) \quad (2)$$

Total reward:

$$R(T) = \sum_{t=1}^T \gamma^t r_t \quad (3)$$

Average reward:

$$R(T) = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T r_t \quad (4)$$

Do you care about future as much as now (and past)?

$\gamma \in [0, 1]$:

- $\gamma \rightarrow 0$: We only care about the current reward not what we'll receive in future
- $\gamma \rightarrow 1$: We care about all rewards equally

$$\mathcal{S} = \{s_0, s_1, s_2\},$$

$$\mathcal{A} = \{a_0, a_1\},$$

$$\mathcal{P}^{a_0} = \begin{bmatrix} 0.5 & 0 & 0.5 \\ 0.7 & 0.1 & 0.2 \\ 0.4 & 0 & 0.6 \end{bmatrix},$$

$$\mathcal{P}^{a_1} = \begin{bmatrix} 0 & 0 & 1 \\ 0 & 0.95 & 0.05 \\ 0.3 & 0.3 & 0.4 \end{bmatrix}.$$




Photo Credit: © https://en.wikipedia.org/wiki/Markov_decision_process

- Policy: The agent's decision
 - Deterministic policy $a = \pi(s)$
 - stochastic policy $\pi(a|s) = P[a_t = a | s_t = s]$
- Value function: how good the agent does in a state

$$V(s) = \mathbf{E} \left[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s \right]$$

- Model: The agent's interpretation of the environment

Not all components are necessary!

Email your questions to

farnaz.adib.yaghmaie@liu.se