

# Policy Gradient



Farnaz Adib Yaghmaie

Linköping University, Sweden  
*farnaz.adib.yaghmaie@liu.se*

March 12, 2021

# What is Policy Gradient?

The most ambitious way of solving an RL problem

- Policy: The agent's decision
- Value function: how good the agent does in a state
- Model: The agent's interpretation of the environment

*Agent's goal:* To learn the policy by directly optimizing the total reward.

$$J = \mathbf{E}_{\tau \sim \pi_{\theta}}[R(T)]$$

How?

Optimization by perturbation:

- Consider a stochastic parametric policy with parameter  $\theta$
- Observe the total reward as a result of perturbation
- Optimize the parameters of the policy by finding  $\nabla_{\theta} J$

# Why to consider a stochastic policy instead of a deterministic one?

- To enable learning by deviating from the deterministic policy
- If a deterministic policy is considered, the agents remains in a local optimum forever.

A simple math rule based on  $\nabla_p \log p = \frac{1}{p}$

$$\nabla_\theta \log p = \nabla_p \log p \nabla_\theta p = \frac{1}{p} \nabla_\theta p.$$

$$\boxed{\nabla_\theta p = p \nabla_\theta \log p} \tag{1}$$

We will have a closer look at the following components in  $J$

$$J = \mathbf{E}_{\tau \sim \pi_{\theta}}[\mathbf{R}(T)]$$

- $\mathbf{R}(T)$ : The total reward
- $\pi_{\theta}$ : The parametric pdf of the policy
- $\tau$ : A sampled trajectory and the expectation is defined over the probability of the trajectory

# How to define the pdf?

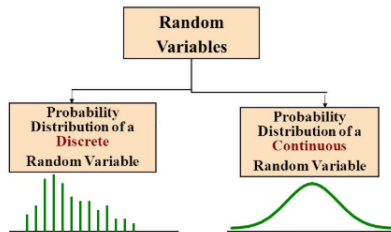
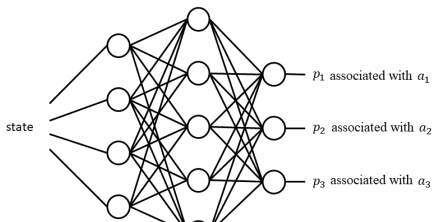


Photo Credit: @ <https://towardsdatascience.com/probability-distributions-discrete-and-continuous-7a94ede66dc0>

Generates the pdf  $\pi_{\theta} = \text{network}(s)$



```
network = keras.Sequential([  
    keras.layers.Dense(30, input_dim=n_s, activation='relu'),  
    keras.layers.Dense(30, activation='relu'),  
    keras.layers.Dense(n_a, activation='softmax')])
```



## Continuous action space

Select a Gaussian distribution as the pdf  $\pi_\theta$  and generate the mean

$$\pi_\theta = \frac{1}{\sqrt{(2\pi\sigma^2)^{n_a}}} \exp\left[-\frac{1}{2\sigma^2}(a - \mu_\theta(s))^T(a - \mu_\theta(s))\right]$$

For example, for a linear policy

$$\mu_\theta(s) = \theta s$$

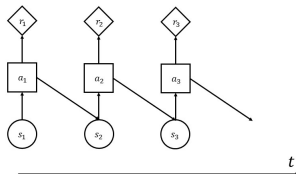
**Discrete:**

Parameterize the pdf

**Continuous:**

Parameterize the mean of a  
Gaussian pdf

## Sampling a trajectory



Select  $a_t \sim \pi_\theta$ ,  $t = 1, \dots, T$  and step the environment

$$\tau = (s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_{T+1})$$

## The probability of the trajectory

$$P(\tau|\theta) = \prod_{t=1}^T p(s_{t+1}|s_t, a_t)p(a_t|\theta).$$

- $p(s_{t+1}|s_t, a_t)$ : the model of the environment
- $p(a_t|\theta)$ : The pdf  $\pi_\theta$  evaluated at  $a_t$ .
  - Discrete action space:  $\pi_\theta = \text{network}(s)$ . So,  $p(a_t|\theta)$  is obtained by indexing into the output vector  $\text{network}(\text{state})$ .
  - Continuous action space:

$$p(a_t|\theta) = \frac{1}{\sqrt{(2\pi\sigma^2)^{n_a}}} \exp\left[-\frac{1}{2\sigma^2}(a_t - \mu_\theta(s_t))^T(a_t - \mu_\theta(s_t))\right]$$

# Love math? Dive in!

$$\begin{aligned}
 \nabla_{\theta} J &= \nabla_{\theta} \mathbf{E}[R(T)] \\
 &= \int_{\tau} \nabla_{\theta} P(\tau|\theta) R(T) \\
 &= \int_{\tau} P(\tau|\theta) \nabla_{\theta} \log P(\tau|\theta) R(T), \quad \text{using log-derivative trick} \\
 &= \mathbf{E}_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \log P(\tau|\theta) R(T)] \\
 &= \mathbf{E}_{\tau \sim \pi_{\theta}} [(\nabla_{\theta} \sum_{t=1}^T \log \underbrace{p(s_{t+1}|s_t, a_t)}_{\text{Dynamics}} + \nabla_{\theta} \sum_{t=1}^T \log p(a_t|\theta)) R(T)] \\
 &= \mathbf{E}_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \sum_{t=1}^T \log p(a_t|\theta) R(T)]
 \end{aligned} \tag{2}$$

But how to compute?

## Similar to the classification task in ML.

Policy Gradient:

$$J = \sum_{t=1}^T R(T) \log p(a_t | \theta)$$

- Number of actions  $n_a$
- Trajectory length  $T$
- Weight  $R(T)$
- State  $s_t$
- Target label for state  $s_t$  and action  $a$
- $p(a_t | \theta)$ : probability of  $a_t$  by the network at

Classification:

$$J_{wcec} = -\frac{1}{M} \sum_{m=1}^M \sum_{c=1}^C w_c y_m^c \log h_{\theta}(x_m, c)$$

- Number of classes  $C$
- Data length  $M$
- Weight  $w_c$
- Image  $x_m$
- Target label for  $x_m$  for class  $c$ ;  $y_m^c$
- $h_{\theta}(x_m, c)$ : The probability of  $x_m$  belongs to class  $c$  by the

# Summary of optimizing the parameters in the discrete action space

PG is similar to the classification task!

- The network should produce probability
- The cost to be optimized is a weighted cross entropy cost
- The weights are  $R(T)$

We can compute  $\mathbf{E}_{\tau \sim \pi_\theta} [\nabla_\theta \sum_{t=1}^T \log p(a_t | \theta) R(T)]$  easily!

$$\nabla_\theta J = \frac{1}{\sigma^2 |\mathcal{D}|} \sum_{\tau \in \mathcal{D}} \sum_{t=1}^T (a_t - \mu_\theta(s_t)) \frac{d\mu_\theta(s_t)}{d\theta}^\dagger R(T).$$

If we consider a linear policy  $\mu_\theta(s_t) = \theta s_t$

$$\nabla_\theta J = \frac{1}{\sigma^2 |\mathcal{D}|} \sum_{\tau \in \mathcal{D}} \sum_{t=1}^T (a_t - \theta s_t) s_t^\dagger R(T).$$



**Discrete:**

Assign a cross entropy cost function and let the ML library optimize the parameter!

**Continuous:**

Use

$$\nabla_{\theta} J = \frac{1}{\sigma^2 |\mathcal{D}|} \sum_{\tau \in \mathcal{D}} \sum_{t=1}^T (a_t - \theta s_t) s_t^{\dagger} R(T)$$

and optimize  $\theta$  with a gradient algorithm, e.g.

$$\theta = \theta + \alpha \nabla_{\theta} J.$$

# Putting all together

We build/consider a parametric pdf  $\pi_{\theta}(s)$ . Then, we iterate:

## 1 Collect data

- Observe  $s$  and sample  $a \sim \pi_{\theta}(s)$ .
- Apply  $a$  and observe  $r$ .
- Add  $s$ ,  $a$ ,  $r$  to the history.

## 2 Update the parameter $\theta$

- We calculate the total reward.
- We optimize the policy by a gradient algorithm.

## PG: The most ambitious way of solving an RL problem

- Directly optimizes the reward for MDP
- No model, no bellman equation
- Random search is a special case
- Can be extremely good or bad
- Take a look at implementation on my github

`Crash_course_on_RL/pg_notebook.ipynb`

# Email your questions to

*farnaz.adib.yaghmaie@liu.se*