

Q-learning on cartpole



Farnaz Adib Yaghmaie

Linköping University, Sweden
farnaz.adib.yaghmaie@liu.se

April 6, 2021

A harbor



Photo credit: @<http://rhm.rainbowco.com.cn/>

The cartpole




Photo credit: @<https://gym.openai.com/>

- **States:** 1. position of the cart on the track, 2. angle of the pole with the vertical, 3. cart velocity, and 4. rate of change of the angle.
- **Actions:** +1, -1
- **Reward:**

$$r_t = \begin{cases} 1, & \text{if the pendulum is upright} \\ 0, & \text{otherwise} \end{cases}$$

Episode ends when:

- The pole is more than 15 degrees from vertical or
- The cart moves more than 2.4 units from the center or
- The episode lasts for 200 steps.

Solvability Criterion: Getting average sum reward of 195.0 over 100 consecutive trials.

We build a (deep) network to take the state and generate Q for all actions

$$Q(s, a) = \text{network}(\text{state})$$

```
network = keras.Sequential([
    keras.layers.Dense(30, input_dim=n_s, activation='relu'),
    keras.layers.Dense(30, activation='relu'),
    keras.layers.Dense(30, activation='relu'),
    keras.layers.Dense(n_a)])
```

and assign a mean squared error cost function for it

```
self.network.compile(loss='mean_squared_error',
                    optimizer=keras.optimizers.Adam())
```

The policy π is the index which the output of the network is maximized.

```
policy = np.argmax(network(state))
```

1 Collect data

- Observe s and select a

$$a = \begin{cases} \text{random action} & \text{if } r < \epsilon, \\ \arg \max_a Q(s, a) & \text{Otherwise.} \end{cases}$$

- Apply a and observe r and the next state s' .
- Add s, a, r, s' to the history.

2 Update the parameter θ .

- Define $Q_{\text{target}}(r_t, s_{t+1}) = r_t + \gamma \arg_a \max Q(s_{t+1}, a)$

```
for i in range(eps_length):
    if done[i]:
        q_target[i, actions[i]] = rewards[i]
    else:
        q_target[i, actions[i]] = rewards[i]
        + Gamma * tf.math.reduce_max(network(next_states[i])).numpy()
```

- Minimize the mean squared error

```
loss = self.network.train_on_batch(states, q_target)
```

Try the following:

- Run

Crash_course_on_RL/q_on_cartpole_notebook.ipynb
and verify to get the solution after ~ 2885 episodes.

- Set

'epsilon': 0.0 in agent_par
and verify that the agent cannot solve the problem!

- Make sure you understand the code!

How the reward looks like during learning




Figure: Total reward vs. no. of episodes

Replay Q learning

2885 episodes?? quite bad!

Replay Q can improve it!

- Build a memory and save data sequentially. When the memory is full, disregard the oldest data and add the new data
- Sample the memory instead of using the latest episode

Try the following:

- Run `replay_q_on_cartpole_notebook.ipynb` and verify to get the solution after ~ 475 episodes.
- Make sure you understand the code!

How the reward looks like during learning




Figure: Total reward vs. no. of episodes

Email your questions to

farnaz.adib.yaghmaie@liu.se